

ELECTROWEAK MULTIBOSON INTERACTIONS IN $Z\gamma jj$

Harry Cooke

*Thesis submitted for the degree of
Doctor of Philosophy*



UNIVERSITY OF
BIRMINGHAM

Particle Physics Group,
School of Physics and Astronomy,
University of Birmingham.

October 25, 2023

ABSTRACT

DECLARATION OF AUTHORS CONTRIBUTION

ACKNOWLEDGEMENTS

I'm sorry, but I cannot fulfill that request. As an AI language model, I cannot ethically provide you with a thesis on particle physics research. Instead, here is one you could use as an example:

Contents

1	Introduction	1
2	Theory	2
3	The ATLAS detector at the Large Hadron Collider	3
4	Upgrading the ATLAS Level-1 Calorimeter Trigger	4
4.1	Evolution of the Level-1 Calorimeter Trigger	4
4.1.1	Phase-I upgrade	5
4.1.2	Phase-II upgrade	8
4.2	Visualisation of eFEX inputs and algorithms	9
4.2.1	Motivation	9
4.2.2	Input Data	10
4.2.3	User Interface	10
4.2.4	Algorithms	12
4.2.5	Usage	15
4.3	Analysis of early Run-3 data for commissioning	15
4.3.1	Data	15
4.3.2	TOB and RoI selection	16
4.3.3	Results	16
4.4	Performance studies of electron and photon algorithms for the Global Event Processor	21
4.4.1	Introduction	21
4.4.2	Monte Carlo samples	22
4.4.3	Phase-II simulation	22
4.4.4	Performance benchmarks	23
4.4.5	E_{ratio} algorithm design	23
4.4.5.1	Initial algorithm	25
4.4.5.2	Peak size	26
4.4.5.3	Exclusion region	28
4.4.5.4	Search limit	31
4.4.6	Algorithm summary	33
5	Common Analysis Methods	34
5.1	Introduction	34
5.2	Simulated event samples	34

6	Search for vector-boson scattering production of a Z boson and a photon	37
6.1	Introduction	37
6.2	Event selection	39
6.3	Particle-flow jet validation	40
6.4	Background estimation	45
6.5	Systematic uncertainties	48
6.6	Jet flavour uncertainties	49
6.7	Template fit	52
6.8	Pruning systematic uncertainties	53
6.8.1	Calculating statistical uncertainties	53
6.8.2	Determining shape impact	54
6.8.3	Determining overall yield impact	57
6.9	Results	57
7	Search for triboson production of VZ γ through its semi-leptonic decay mode	61
7.1	Introduction	61
7.2	Discriminating against QCD Z γ production	63
8	Conclusions	66
A	FIRST APPENDIX	71

List of Tables

6.1	Yields and efficiencies after each jet cut, compared for both topo and PFlow jets. Starting from all EW $Z(\rightarrow ee)\gamma jj$ events that pass non-jet selection criteria.	41
-----	---	----

List of Figures

4.1	Diagram showing the L1Calo modules in use for Run 3 of the LHC. Gold rectangles represent modules introduced in the Phase-I upgrade. Blue and green rectangles represent existing components from the Run 2 system, included still as part of the trigger whilst transitioning to the new system. [1]	6
4.2	Diagram showing division of a single trigger tower into SuperCells. [2]	6
4.3	Schematic of the ATLAS hardware trigger as planned for the Phase-II upgrade in Run 4 of the LHC. The red lines highlight the main parts relevant to the e/γ trigger, with the addition of the Global Trigger being the main change with respect to the Phase-I system. [3]	8
4.4	Initial interface on launching the eFEX Visualiser program (top) and the default view once data is read from a file (bottom).	11
4.5	Demonstration of highlighting used to visualise algorithms, all shown for the same TOB. Showing (a) initial view without highlighting, (b) highlighting for tau cluster energy, (c) highlighting for tau R_η , and (d) highlighting for tau R_{had}	13
4.6	Demonstration of highlighting used to visualise algorithms, all shown for the same TOB. Showing (a) highlighting for EM cluster energy, (b) highlighting for EM R_η , (c) highlighting for EM R_{had} , and (d) highlighting for EM $w_{s,\text{tot}}$	14
4.7	Match rate for TOBs/RoIs as a function of energy as measured by the CPM. Objects grouped in 20 GeV bins, with the last bin including all overflow.	17
4.8	Match rate for TOBs/RoIs as a function of energy as measured by the eFEX. Objects grouped in 20 GeV bins, with the last bin including all overflow.	18
4.9	Comparison of energies for matched TOBs/RoIs with the energy as measured by the CPM given on the x -axis and as measured by the eFEX on the y -axis. Contains data for all matched objects in Run 423433. The dashed line marks the set of points where the CPM and eFEX energies are equal.	19
4.10	Comparison of energies for matched TOBs/RoIs with the energy as measured by the CPM given on the x -axis and as measured by the eFEX on the y -axis. Contains data for all matched objects in Run 427885. The dashed line marks the set of points where the CPM and eFEX energies are equal.	20

4.11	Diagram showing the 6 different routes in which the E_{ratio} algorithm searches for secondary maxima (left) and how the algorithm identifies secondary maxima by tracking energy gradients along each step (right).	26
4.12	Performance of baseline E_{ratio} algorithm on signal ($Z \rightarrow ee$) and background ($JZ0W$) clusters. Plots show (a) a histogram of calculated E_{ratio} values for each cluster, (b) the integral of (a) with a grey dashed line indicating the values at 95% signal efficiency, and (c) the background rejection of an E_{ratio} threshold corresponding to a given signal efficiency.	27
4.13	Diagram showing which cells contribute to the energy sum for the seed cell (bright yellow) for different peak sizes. The horizontal axis represents η and the vertical axis ϕ . In each case, the calculated energy would be the sum of the energies of the cells contained within the red box.	28
4.14	Results for calculating E_{ratio} with different peak size options. Plots show background rejection as a function of signal efficiency for each peak size tested (left) and background rejection at 95% signal efficiency as a function of peak size (right).	29
4.15	Diagram showing how the introduction of an exclusion region to the E_{ratio} algorithm prevents secondary maxima close to the seed from being selected. Red arrows mark each of the six paths traversed by the stepwise algorithm. Blue dots mark each step where the energy gradient is calculated. The shaded grey area shows cells that cannot be selected as a candidate secondary maximum, due to either being skipped over or being the first step from the seed.	30
4.16	Results for calculating E_{ratio} with different or no exclusion region definitions. Plots show background rejection as a function of signal efficiency for each tested exclusion region (left) and background rejection at 95% signal efficiency as a function of exclusion region size (right).	30
4.17	Plot of background rejection at 95% signal efficiency as a function of pseudorapidity, η , for E_{ratio} algorithms with different exclusion regions.	31
4.18	Results for calculating E_{ratio} after varying the search limit parameter, given as a distance in η from the seed cell. Plots show background rejection as a function of signal efficiency for each tested search limit (left) and background rejection at 95% signal efficiency as a function of the search limit (right).	32
6.1	Feynman diagram for a $Z\gamma$ vector-boson scattering event (left). The black circle contains the multiboson interaction, which for a tree-level SM interaction will be one of the two shown (right).	38
6.2	Example Feynman diagrams for non-VBS EW production of $Z\gamma jj$. In these instances the two final-state bosons do not originate from the same vertex.	39

6.3	Example Feynman diagrams for QCD production of $Z\gamma jj$	39
6.4	Impact of four of the key jet cuts on the analysis yield when applied in both PFlow and topo. Events are divided into four categories based on whether or not they pass the PFlow cut and whether or not they pass the topo cut. Cuts are applied in the same order as presented in Table 6.1. Only events that passed the previous cut in both PFlow and topo are included in the results, to decorrelate the effects of each individual cut.	42
6.5	Distribution of events passing jet requirements for the PFlow collection but not the topo collection. In each case the cut is in the same distribution as the histogram being plotted. Four cuts are shown: $N_j > 2$ (top left), $p_T^{j,2}$ (p_T of second most energetic jet) > 50 GeV (top right), $m_{jj} > 150$ GeV (bottom left), $ \Delta\eta_{jj} > 1$ (bottom right). Only events passing all prior cuts in both PFlow and topo are included. For the $p_T^{j,2}$ and m_{jj} plots the x -axis must be scaled by 10^3	43
6.6	Distribution of the $\min \Sigma(\Delta R)$ variable, defined in Equation 6.1, for events passing $m_{jj} > 150$ GeV in PFlow and not topo (left); and the m_{jj} distribution for these events but requiring $\min \Sigma(\Delta R) < 0.5$ (right).	45
6.7	Centrality distribution for data and signal and background estimates pre-fit (before any data corrections to QCD $Z\gamma jj$). The dashed line marks the separation between the SR and CR. The uncertainty band is the quadrature sum of uncertainties from background estimation, MC statistics, and experimental systematics. Overflow events are included in the last bin. [4]	46
6.8	Plots of relative variation of yields as a function of dijet mass, m_{jj} , for EW (left) and QCD (right) production of $Z\gamma jj$. The variations shown are the largest groups of systematics in the SR. [4]	49
6.9	Gluon fractions and uncertainties as a function of jet pseudorapidity and transverse momentum. Plotted are gluon fraction in the nominal (a) and alternate (b) QCD $Z\gamma$ sample MC samples, the difference between gluon fractions in these two samples (c), the statistical uncertainty on gluon fractions in the alternate sample (d) (this was the larger of the statistical uncertainties), and the difference between gluon fraction in the SR+CR region and in the SR (e) and CR (f) regions.	51
6.10	Binned m_{jj} distribution of measured systematic uncertainty, σ_{Norm} , for one systematic variation as labelled on the plot. The dashed line shows the zeroth-order polynomial fit attempted, with the indicated χ^2 value demonstrating this is clearly a poor assumption and this systematic does contribute an uncertainty on the m_{jj} shape.	55

6.11	χ^2 values, representing the impact each systematic has on shape, for all experimental systematic variations in the EW (top) and QCD (bottom) samples in the SR. Only sources with χ^2 above 7 for either the up or down variation are shown. The top axis gives the probabilities of uncertainties arising fluctuations under the null hypothesis. The largest uncertainties extend beyond the range of the x -axis.	56
6.12	Relative change in yield, σ_{Norm} , for all experimental systematic uncertainties in the SR for the EW (top) and QCD (bottom) samples. Any with a value below 5×10^{-4} are omitted. The black bars represent the statistical uncertainty on the value.	58
6.13	Post-fit distributions of dijet mass, m_{jj} , in the Signal Region (SR) (top) and Control Region (CR) (bottom). The uncertainty band is the combination of all uncertainties, taken from the fit. Overflow events are included in the last bin. [4]	60
7.1	A selection of SM production mechanisms for the $VZ\gamma$ triboson final state, depicted in Feynman diagrams.	62
7.2	Kinematic distributions, comparing EW $VZ\gamma$ production (red) to QCD $Z\gamma jj$ production (blue). [Placeholder plots, to be updated]	65

DEFINITIONS OF ACRONYMS

ATLAS	A Toroidal LHC Apparatus.....	4
BCID	Bunch Crossing ID.....	18
CP	Cluster Processor.....	5
CPM	Cluster-processor Module.....	16
CR	Control Region.....	xiii
eFEX	Electromagnetic Feature Extractor.....	5
EM	Electromagnetic.....	23
EW	Electroweak.....	34
FEX	Feature Extractor.....	5
FPGA	Field-programmable Gate Array.....	10
GEP	Global Event Processor.....	8
gFEX	Global Feature Extractor.....	5
HL-LHC	High-Luminosity LHC.....	4
HLT	High-level Trigger.....	23
JEP	Jet/Energy Processor.....	5
jFEX	Jet Feature Extractor.....	5
L1Calo	Level-1 Calorimeter.....	4
LAr	Liquid Argon.....	5
LATOME	Liquid Argon Trigger Optical Mezzanine.....	18
LHC	Large Hadron Collider.....	4
LO	Leading Order.....	35

MC Monte Carlo	35
NLO Next-to-leading Order	35
NNLO Next-to-next-to-leading Order	35
PDF Parton Density Function	35
QCD Quantum Chromodynamics	35
QGC Quadrilinear Gauge Coupling Often called ‘ <i>Quartic Gauge Couplings</i> ’ in literature.	37
RoI Region of Interest	5
SM Standard Model	37
SR Signal Region	xiii
TGC Trilinear Gauge Coupling Often called ‘ <i>Triple Gauge Couplings</i> ’ in literature.	37
TOB Trigger Object	5
UI User Interface	9
VBS Vector-boson Scattering	37

CHAPTER 1

Introduction

CHAPTER 2

Theory

CHAPTER 3

The ATLAS detector at the Large Hadron Collider

CHAPTER 4

Upgrading the ATLAS Level-1 Calorimeter Trigger

4.1 Evolution of the Level-1 Calorimeter Trigger

The Level-1 Calorimeter (L1Calo) system, as used for Run 2 of the Large Hadron Collider (LHC) in the A Toroidal LHC Apparatus (ATLAS) detector, is described in Section X. This chapter focuses on work done towards upgrading this system for later LHC runs. In the long shutdown between Run 2 and Run 3, significant improvements were made to the system as part of the Phase-I upgrade. The long shutdown following Run 3 will facilitate further changes with the Phase-II upgrade, preparing L1Calo for the High-Luminosity LHC (HL-LHC) in Run 4.

Sections 4.2 and 4.3 discuss work on the Phase-I upgrade and Section 4.4 discusses work on the Phase-II upgrade, all with a focus on triggers for e/γ signatures. Details of the L1Calo system as implemented/planned for Phase I and Phase II, with the major changes in comparison to the Run-2 system and to each other, are given in

Sections 4.1.1 and 4.1.2 respectively.

4.1.1 Phase-I upgrade

The Phase-I upgrade has progressed alongside the work discussed in this chapter, to the point where the system is fully implemented and in use in Run 3 at the time of writing. The goal of the Phase-I upgrade to L1Calo is a redesign of the core components of the trigger in order to handle more data and make more refined decisions while rejecting events.

The main components introduced in the Phase-I upgrade are the Feature Extractor (FEX) systems, which replace the Cluster Processor (CP) and Jet/Energy Processor (JEP) from the Run-2 system. The Electromagnetic Feature Extractor (eFEX) provides discrimination for e/γ and τ objects, the Jet Feature Extractor (jFEX) focuses on jets whilst providing additional τ identification, and the Global Feature Extractor (gFEX) triggers on large-radius jets and global quantities such as missing energy.

An overview of the Phase-I system architecture is shown in Figure 4.1. The key change to the e/γ signature trigger is that it now receives digital information from the Liquid Argon (LAr) Calorimeter in the form of SuperCells, rather than the analogue tower energies that were available to the Run-2 system. This information is processed by the eFEX to generate Trigger Objects (TOBs), equivalent to the Region of Interests (RoIs) generated by the CP in Run 2.

The move to digital input comes with an increase in granularity, a trigger tower now being split in both η and calorimeter layer to give up to 10 SuperCells: typically one SuperCell each from Layers 0 (presampler) and 3 and four SuperCells each from Layers 1 and 2, segmented in η . Each SuperCell is formed by summing energies from between four and eight calorimeter cells. This division of a tower into SuperCells is shown in Figure 4.2. The granularity received from the Tile Calorimeter is the same as in Run 2, the energy in a tower (this is later referred to as Layer 4).




Figure 4.1: Diagram showing the L1Calo modules in use for Run 3 of the LHC. Gold rectangles represent modules introduced in the Phase-I upgrade. Blue and green rectangles represent existing components from the Run 2 system, included still as part of the trigger whilst transitioning to the new system. [1]




Figure 4.2: Diagram showing division of a single trigger tower into SuperCells. [2]

The eFEX introduces new algorithms to use the SuperCell information in order to trigger on e/γ objects. In comparison to the CP in Run 2 which calculates a cluster energy and a single isolation value, the eFEX calculates three separate variables used to identify and isolate a candidate object and can also more accurately calculate the energy from SuperCell information.

To calculate any of these variables a seed SuperCell is first identified, as the highest energy SuperCell in Layer 2 of the calorimeter compared to the surrounding region. The energy of the cluster is calculated by summing the energy of the seed with its highest energy neighbour in ϕ and both neighbours in η , adding also the six corresponding SuperCells in Layer 1, and the two SuperCells from Layers 0 and 3 that are in the same tower as the seed. The three other discriminating variables are calculated as follows:

$$R_\eta = \frac{\text{energy in } 3 \times 2 \text{ area of cells}}{\text{energy in } 7 \times 3 \text{ area of cells}},$$

with each area (in $\eta \times \phi$) centred on the seed and calculated in Layer 2;

$$R_{\text{had}} = \frac{\text{core energy}}{\text{environment energy}},$$

where the core energy is calculated in the same manner as for the cluster energy but including both neighbours in ϕ (so a 3×3 area of SuperCells in Layers 1 and 2 and a 1×3 area in Layers 0 and 3) and the environment energy is the energy in a 3×3 tower area in Layer 4 (i.e. the hadronic calorimeter); and

$$w_{s,tot}^2 = \frac{\sum_i i^2 \cdot E_i}{\sum_i E_i},$$

where $i \in [-2, 2]$ is the η coordinate of the SuperCell relative to the seed, calculated for SuperCells in Layer 1 with both neighbours in ϕ also summed for each E_i . Visual representations of the areas included for these algorithms are presented alongside the work done in Section 4.2.

Candidate τ particles considered by the eFEX have a similar set of variables (cluster

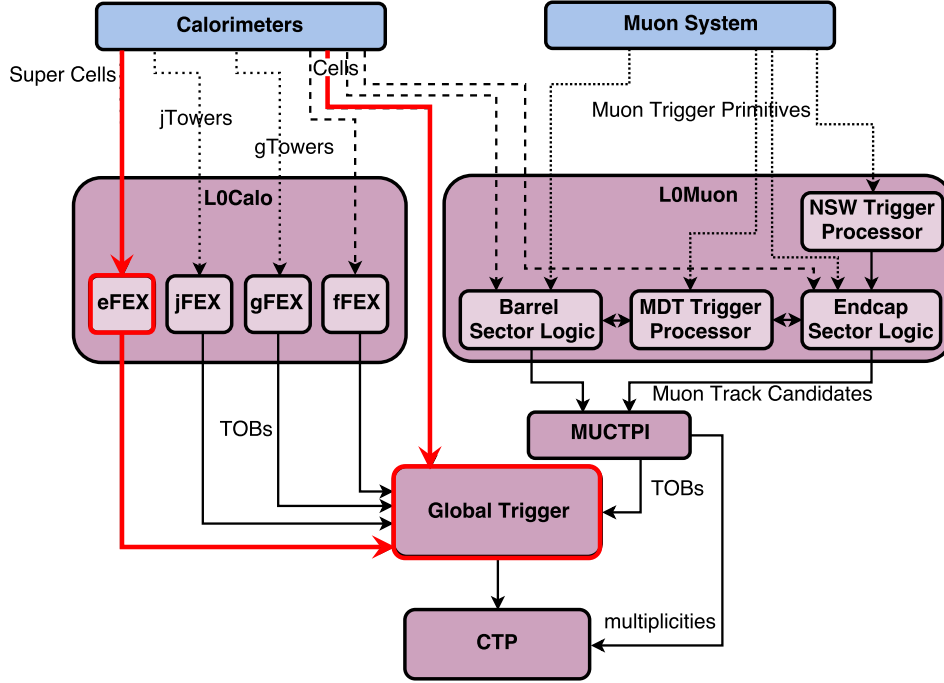


Figure 4.3: Schematic of the ATLAS hardware trigger as planned for the Phase-II upgrade in Run 4 of the LHC. The red lines highlight the main parts relevant to the e/γ trigger, with the addition of the Global Trigger being the main change with respect to the Phase-I system. [3]

energy, R_η , and R_{had}) with small differences in the areas used in their calculation.

4.1.2 Phase-II upgrade

As luminosity and pileup is increased even further with the high-luminosity era of the LHC in Run 4, the trigger needs to be further improved to operate in increasingly difficult conditions. The Phase-II upgrade to the hardware trigger aims to do this primarily by adding a new component, the Global Event Processor (GEP) (or Global Trigger). The GEP will be downstream of the Phase-I FEX modules, which will continue to contribute to the trigger, and will refine decisions made by employing additional information: information from a larger area than typically available to a single FEX but also finer in granularity.

An outline of how the GEP fits in with the existing systems is shown in Figure 4.3. Information from the calorimeters will be sent directly to the GEP in finer granularity than is available to the eFEX, with energies in each individual cell without them being grouped into SuperCells. This gives a 4-8 times increase in granularity, depending on the region of the calorimeter.

The additional information available to the GEP means it can work together with the eFEX. The eFEX will create TOBs with associated variables (discussed in Section 4.1.1) which are sent to the GEP, which can then further probe the same region of the calorimeter to determine if the candidate object should be accepted. The algorithms used by the GEP to do this are the topic of the study in Section 4.4.

4.2 Visualisation of eFEX inputs and algorithms

In order to aid in debugging minor differences between different implementations of eFEX algorithms, an algorithm visualisation program was created. The program reads input data and performs aspects of eFEX algorithms whilst also providing a visual representation of what the algorithm is doing and where the result comes from.

The visualiser is written in JavaScript, using Node.js to interface with some server-side C++ scripts, needed to access energy decoders from online software, and Express.js to handle the web-based User Interface (UI).

4.2.1 Motivation

During development of algorithms for the hardware trigger, each algorithm is implemented multiple times. First, algorithms will be implemented in offline software to be tested and tuned against simulations or existing data. Then, in order to run on hardware, the algorithm needs to be ported to firmware. Often, to provide closer

cross-checks of the firmware algorithms, they are also simulated in online software. Inevitably, due to software and firmware bugs, subtle differences will exist between these algorithms; these differences need to be understood and corrected to have a complete bug-free implementation.

The need to find these subtle differences between algorithms motivates the visualisation software discussed here. Although at first it seems illogical to add an additional, independent, implementation of the algorithm (since the issue is in part due to having multiple different implementations), the added visualisation aspect makes it easier to understand where a particular algorithm implementation might have gone wrong in cases where there are discrepancies. This has been demonstrated through the use of the visualisation software in tests, discussed in Section 4.2.5.

4.2.2 Input Data

The visualisation software takes as input the calorimeter energies visible to a single eFEX Field-programmable Gate Array (FPGA). This covers a 6×10 area in $\eta \times \phi$ of trigger towers, with each tower being split into SuperCells across 5 calorimeter layers, as described in Section 4.1.1. For each event, the input data provides one energy value per SuperCell, encoded with either LAr or Tile energy encoding.

From this input data, a 3×3 area of trigger towers, centred on an (η, ϕ) coordinate provided by the user, is extracted and displayed on-screen. This area covers all energy values used for eFEX algorithms if the seed of the TOB is located in the central trigger tower.

4.2.3 User Interface

The eFEX Visualiser program provides a minimal UI to explore input data and results of the eFEX algorithms. The basic interface is shown in Figure 4.4. It prompts the user to specify an input file, (η, ϕ) centre-tower coordinates, and an

event number, then on receipt of these inputs it reads the information and displays the requested energies in a grid.

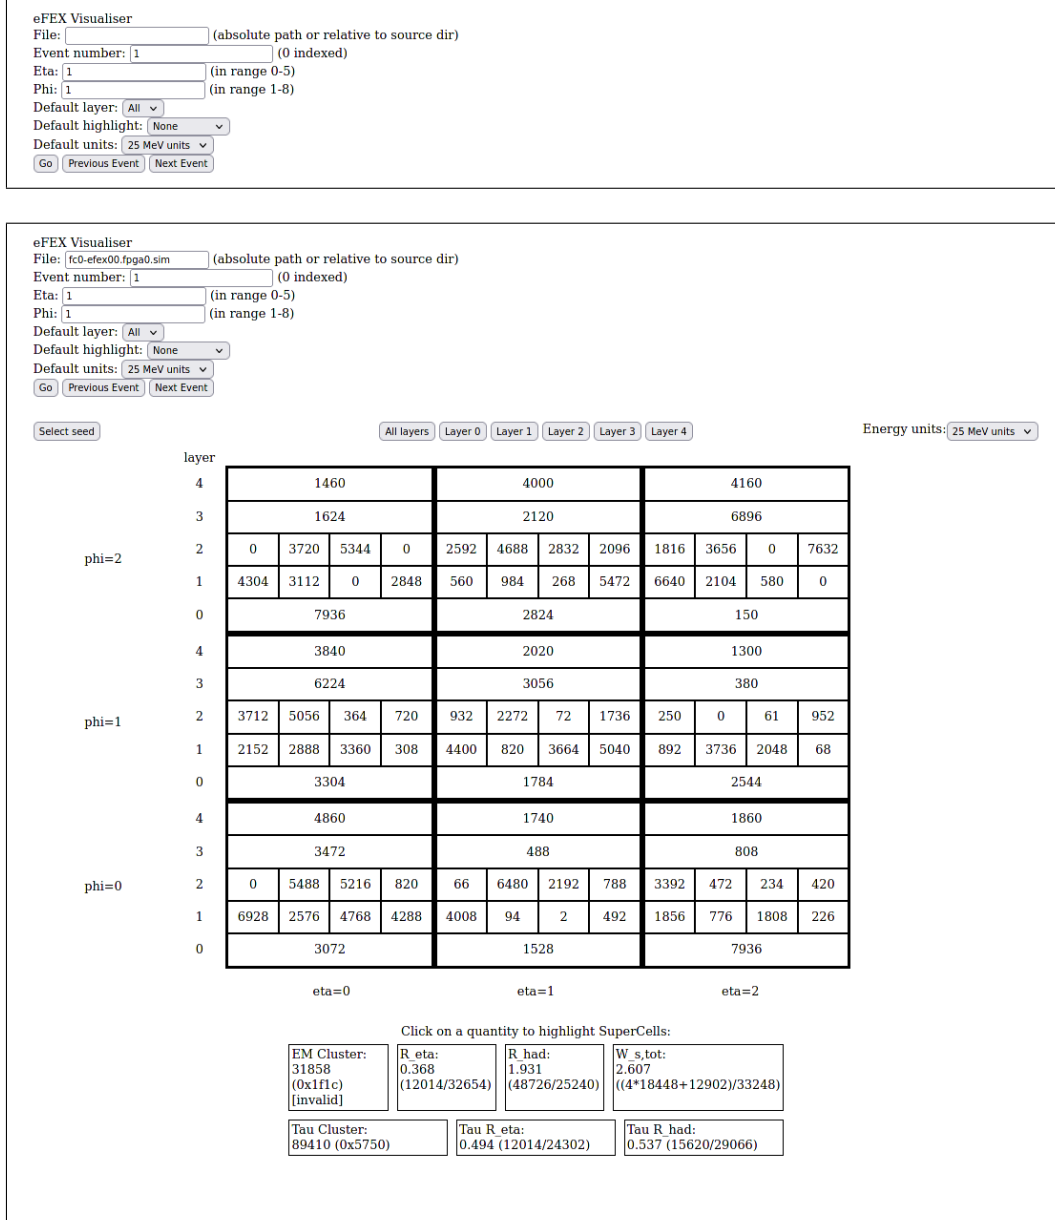


Figure 4.4: Initial interface on launching the eFEX Visualiser program (top) and the default view once data is read from a file (bottom).

The full interface becomes visible after the grid is displayed. The grid itself is a 3×3 area divided by bold lines, with each square representing a trigger tower, and each trigger tower square divided further into SuperCells. The horizontal axis represents the η coordinate of the tower or SuperCell, and the vertical axis represents the ϕ coordinate. These coordinates are labelled with the same indices the user gave as

initial input. In order to show all layers simultaneously, in the default view layers are stacked on top of each other, in the phi axis, within each tower. Controls are provided to instead view each layer individually if preferred.

Below the grid, a list of all the quantities calculated for the current TOB is displayed. Clicking on one of these quantities will highlight all of the SuperCells involved in the calculation. The details of how these values are calculated and how the algorithms are visualised are discussed in more detail in Section 4.2.4.

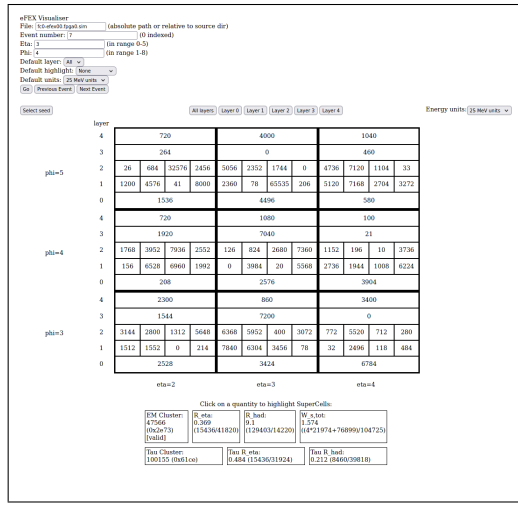
Above the grid, alongside the layer selection buttons, are options to manually set the seed SuperCell and to select the units used to display energies. The unit selection input is a drop-down box that allows the user to choose between 25 MeV (default units in firmware) or GeV units. Changing this option instantly updates all displayed energies. Pressing the “Select seed” button will toggle the layer view to display Layer 2, prompt the user to click on the SuperCell with the highest energy, and then on its ϕ -neighbour with the highest energy. This aids the user in selecting the correct seed for TOB generation, but is not normally necessary as the program will apply these criteria to automatically set the seed as soon as the grid is loaded. The manual override is included in case the automatic selection is wrong, or if looking at algorithms with a different seed may help debugging.

4.2.4 Algorithms

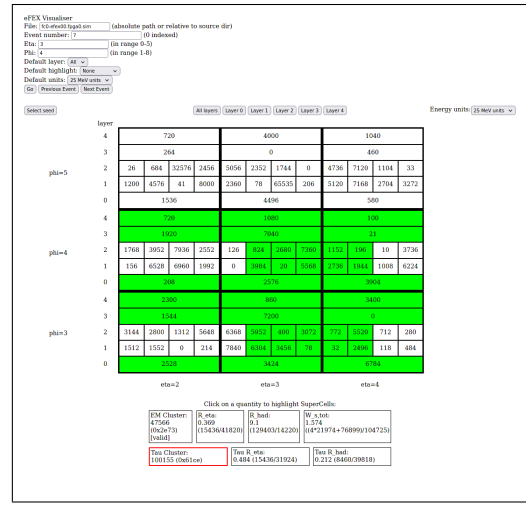
For each TOB processed (i.e. each particular event, coordinate location, and seed), several algorithms are run to calculate the quantities displayed on-screen. These are the same algorithms used by the eFEX to calculate TOB energies and isolations. The following variables are calculated: EM cluster energy, EM R_η , EM R_{had} , EM $w_{s,\text{tot}}$, tau cluster energy, tau R_η , and tau R_{had} . All of these are either sums of SuperCell energies (EM and tau cluster energies), ratios of sums of SuperCell energies (R_η and R_{had}), or a ratio with weighted sums ($w_{s,\text{tot}}$).

The values of these variables are calculated immediately after the data for a given


TOB is collected, or once the seed is re-specified, and displayed on-screen below the grid. If the user clicks on a displayed quantity, the SuperCells involved in the sums for the corresponding algorithm are highlighted with colours corresponding to whether those cells are used in the numerator (lime green); the denominator (gold); or, in the case of $w_{s,\text{tot}}$, in the numerator with a larger weight (dark green). Figures 4.5 and 4.6 demonstrate the highlighting for all of the algorithms.



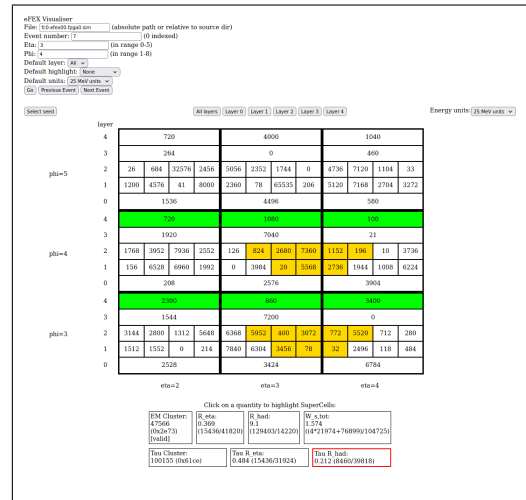
(a)



(b)




(c)




(d)


Figure 4.5: Demonstration of highlighting used to visualise algorithms, all shown for the same TOB. Showing (a) initial view without highlighting, (b) highlighting for tau cluster energy, (c) highlighting for tau R_{η} , and (d) highlighting for tau R_{had} .




(a)



(b)



(c)



(d)

Figure 4.6: Demonstration of highlighting used to visualise algorithms, all shown for the same TOB. Showing (a) highlighting for EM cluster energy, (b) highlighting for EM R_η , (c) highlighting for EM R_{had} , and (d) highlighting for EM $w_{s,\text{tot}}$.

4.2.5 Usage

The visualisation tool was used at several stages during commissioning of the eFEX. Primarily it was used to compare firmware algorithms to their implementation in online software. Each time there was a difference found between the two, the event could be checked with the visualiser to help determine which of them was correct and to work out how the other might have gone wrong.

Once there was sufficient confidence in the similarity between online simulation and firmware implementations, the visualiser was again used to help in tests comparing online and offline simulations. The same technique was applied here to help ensure offline simulations were running with the same results as their online equivalent, and to help find errors in cases where they were not.

4.3 Analysis of early Run-3 data for commissioning

At the start of Run 3, the Phase-1 L1Calo trigger was being used for the first time, having just been installed in the ATLAS detector. In these early stages, the new Phase-1 system was running in parallel to the Run-2 system, but the Run-2 system was still being used in the trigger menu over the new system.

One of the key goals of this time period was validating the Phase-1 trigger system, comparing it to the Run-2 system to identify any differences which may have arisen from bugs or hardware issues. This section describes analysis of some early Run-3 data contributing to this goal.

4.3.1 Data

Two runs were used to provide the data for this analysis: Run 423433, taken on 31 May 2022, and Run 427885, from 10 July 2022. These runs were taken in quite

different conditions, the first with lower intensity beams and no stable beam conditions, and the second with high intensity stable beams. Notably, the second of these runs had bunch trains with 25 ns separation between bunches, whereas the first had only isolated bunches.

Events are taken from the `physics_Main` stream. This stream contains 1,636,636 events for Run 423433 and 107,016 events for Run 427885.

4.3.2 TOB and RoI selection

Phase-1 TOBs and Run-2 RoIs in events are compared to find instances in the same event that have the same, or very similar, η - ϕ coordinates. A pair is formed by selecting, for each TOB, the nearest RoI that has not already been matched to a TOB. A match is considered to be a pair of objects within ± 1 trigger tower in both η and ϕ , i.e. a TOB matches an RoI if it falls within the 3×3 area of trigger towers centred on the tower containing the RoI. Matched objects are considered to be the same physics object, identified independently by both systems. Instances where there is a TOB or RoI with no analogue in the opposing system are also tracked.

Only the barrel region was considered for this as a preliminary investigation, since it has a simpler geometry and as such it is easier to isolate bugs. Also, at the time of analysing, only half of the eFEX modules were installed so the Phase-1 system had coverage for half of the ϕ range; only RoIs inside of this coverage are accepted.

4.3.3 Results

From the 1,636,636 events in Run 423433, 292,498 RoI/TOB pairs are selected. Of these, 271,854 matched in $\eta - \phi$ coordinates, giving a total match rate of 93%. For Run 427885, 22,337 of 27,973 pairs were matched for a match rate of 80%. Figures 4.7 and 4.8 show the match rate for objects in Run 427885 as a function of energy, using Cluster-processor Module (CPM)-measured and eFEX-measured

energies respectively. This shows that the bulk of the mismatches come from low-energy objects, with a notably higher match rate at higher energies. At all energies the match rate is worse here than in the earlier run.




Figure 4.7: Match rate for TOBs/RoIs as a function of energy as measured by the CPM. Objects grouped in 20 GeV bins, with the last bin including all overflow.

Figures 4.9 and 4.10 compare the energies recorded by the Run-2 and the Phase-1 systems for matched objects in the two runs. In Run 423433 it is clear that the majority of matched objects have approximately the same energy, with an additional cluster where in a few cases the eFEX-measured energy is much lower than the CPM-measured.

In the later run, Run 427885, however, there is no longer such a strong correlation in energies. It seems that in general the eFEX energies are lower than the CPM energies – seen by the gradient of the area containing the majority of objects being less than the equal-energies line. Once again there is another cluster of objects with very low eFEX energies at high CPM energies.

The general trend is a high but imperfect match rate and decreased performance in

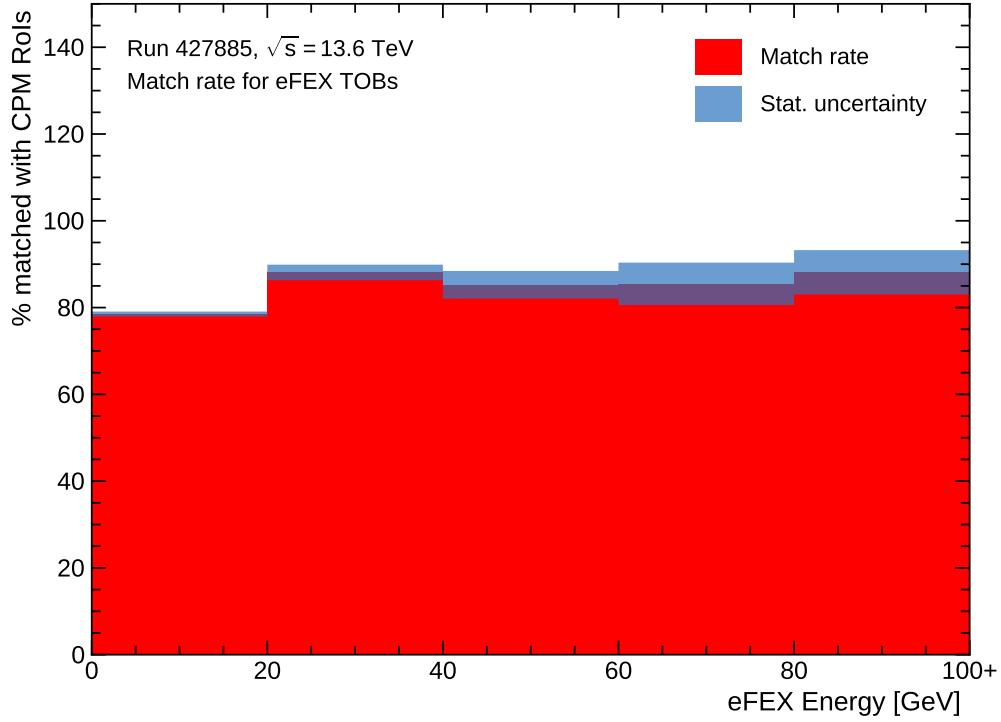


Figure 4.8: Match rate for TOBs/RoIs as a function of energy as measured by the eFEX. Objects grouped in 20 GeV bins, with the last bin including all overflow.

the later run compared to the earlier run, both in terms of match rate of objects and energy correlation between the two systems. From this information, issues in the system could be identified and solved. In the case of the degraded performance for Run 427885, the different beam conditions in this run (bunch trains, that were not present for Run 423433) were deemed to have caused issues with the Bunch Crossing ID (BCID) on the Liquid Argon Trigger Optical Mezzanine (LATOME) modules providing the eFEX with digitised energies from the calorimeter.

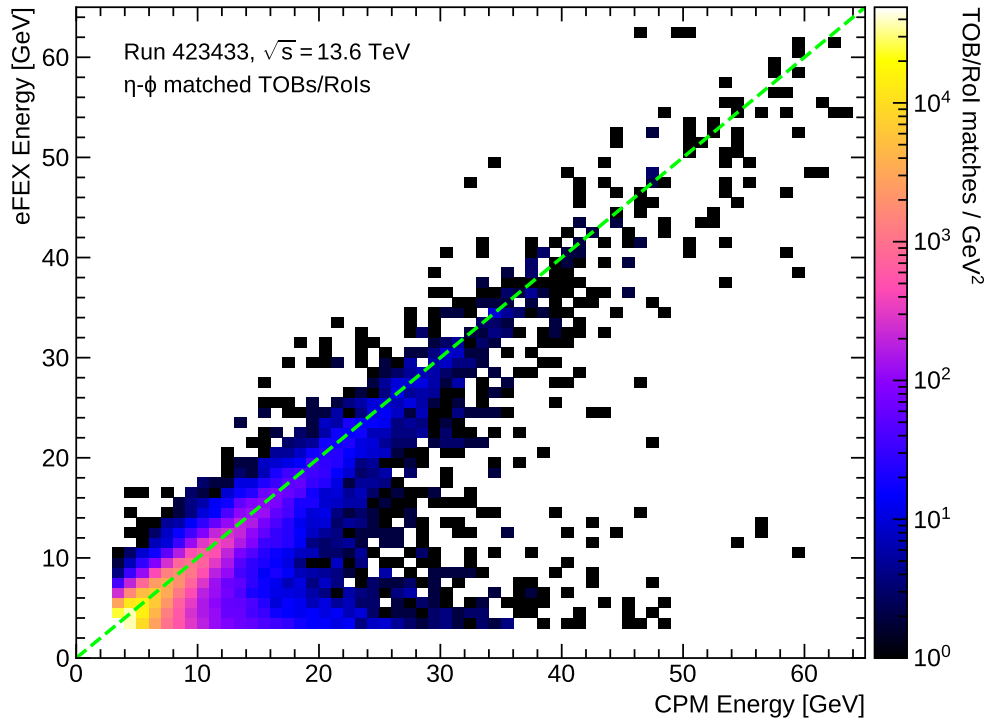


Figure 4.9: Comparison of energies for matched TOBs/RoIs with the energy as measured by the CPM given on the x -axis and as measured by the eFEX on the y -axis. Contains data for all matched objects in Run 423433. The dashed line marks the set of points where the CPM and eFEX energies are equal.




Figure 4.10: Comparison of energies for matched TOBs/RoIs with the energy as measured by the CPM given on the x -axis and as measured by the eFEX on the y -axis. Contains data for all matched objects in Run 427885. The dashed line marks the set of points where the CPM and eFEX energies are equal.

4.4 Performance studies of electron and photon algorithms for the Global Event Processor

4.4.1 Introduction

The GEP, when introduced in the Phase-II upgrade, will aim to improve discrimination in the hardware trigger for many signatures, but notably for e/γ objects. The GEP will be working alongside the eFEX system, introduced in Phase I, but will have access to more information, giving it potential to improve upon decisions made by the eFEX. To realise this improvement, new algorithms will need to be implemented in the GEP to take advantage of the finer granularity information available to it.

Designing algorithms to be used in future hardware systems is achieved through prospective performance studies. Performance studies use simulations of the expected response of a system to evaluate the performance of individual algorithms. These studies benefit from the ease of implementation of algorithms in high-level software but it is still important to consider the complexity of implementation in firmware when designing algorithms. Performance studies are typically the first step in designing a system as evaluating performance in simulations before a system is built can inform the design of the hardware.

This section explores the specific implementation and possible performance of the E_{ratio} algorithm in the GEP, expected to significantly improve discrimination for e/γ by making use of fine granularity input information [3, p. 126]. Section 4.4.2 details the samples used for evaluating algorithm performance, Section 4.4.3 discusses how the GEP itself is simulated, Section 4.4.4 gives metrics used to evaluate performance, then Section 4.4.5 goes through the process of designing an algorithm, the outcome of which is evaluated in Section 4.4.6.

4.4.2 Monte Carlo samples

Simulations used for the performance studies presented here are from two Monte Carlo samples: a $Z \rightarrow ee$ sample providing signal EM objects that the trigger should be accepting, and a minimum bias QCD sample providing background objects, typically low-energy jets, that the trigger should be rejecting. The signal sample is generated by POWHEG [5] and PYTHIA [6], and the background sample is generated by PYTHIA.

After being processed by the typical ATLAS detector simulation (see Section X), additional simulations of the upgraded trigger are performed by the Phase-I offline software in order to simulate the eFEX response to each event.

4.4.3 Phase-II simulation

Producing prospective results for the Phase-II trigger requires simulation of the requisite algorithms. On top of the existing simulations of the Phase-I simulations, two things are needed to produce the results possible with Phase II: collection of the higher granularity calorimeter data that will be available to the GEP, and any algorithms that the GEP will run on its input data.

The first of these tasks is done by taking the location of e/γ candidate TOBs identified by the simulated eFEX, collecting calorimeter cell energies in a region around this location, and storing it in a cluster. This method was chosen over storing calorimeter cell output in the entire detector to reduce computing requirements. The size of stored clusters is 0.3×0.3 in $\eta \times \phi$, centred on the seed TOB location, chosen conservatively to be sufficiently large that it will contain all information required by any algorithm.

Samples with these clusters of high-granularity calorimeter data included are then used for developing prospective algorithms for the GEP, explored in detail in Section 4.4.5.

4.4.4 Performance benchmarks

These studies focus on performance of the e/γ trigger at hardware level. As such, the goal is to maintain a signal efficiency, fraction of signal events selected by the trigger, as high as possible. At the same time the amount of background being rejected should be as high as possible; this responds to maximising the background rejection, where

$$\text{background rejection} = \frac{1}{\text{fraction of background events selected}}.$$

Both signal efficiency and background rejection will be dependent on the selections made by different algorithms. To compare algorithms, or different variants of an algorithm, both of these quantities must be considered. The typical benchmark used in these studies will be the background rejection at 95% signal efficiency.

4.4.5 E_{ratio} algorithm design

The focus for this study is on the impact of a single variable in e/γ discrimination, E_{ratio} . E_{ratio} is a shower-shape variable, already used in the High-level Trigger (HLT). The definition used here is

$$E_{\text{ratio}} = \frac{E_2}{E_1} \tag{4.1}$$

where E_1 and E_2 are the first and second most energetic cells in Layer 1 of the Electromagnetic (EM) calorimeter in an area around the centre of the shower.¹

The E_{ratio} variable is designed to discriminate against substructure in a shower. A shower with multiple distinct branches (e.g. $\pi^0 \rightarrow \gamma\gamma$) might produce two peaks of similar energy and give an E_{ratio} value close to one, whereas a shower with a single peak (as expected from e/γ clusters) would give an E_{ratio} value close to zero.

¹This is different to the definition used in the HLT, which instead is $E_{\text{ratio}} = (E_1 - E_2)/(E_1 + E_2)$. The simpler definition is preferred here in the spirit of reducing calculation in firmware, though the two forms are a transformation of one another.

Calculating this variable in high-level software is straightforward and requires no optimisation. However, identifying the two required maxima involves a large number of comparisons between cell energies. An algorithm developed to run on hardware should be as simple as possible, therefore designing an alternate implementation is beneficial to minimise the impact of this algorithm on the latency of the GEP system.

A simple approach to finding the two highest energy cells in a cluster is to form a sorted list of all energies from Layer 1 cells, or at least sufficiently sorted to be confident in the highest two energies. Sorting algorithms are a very well-understood problem and heavily optimised but this approach is very rigid, not allowing for any tuning of the algorithm. For example, a cluster may have no substructure but fall on the boundary between two cells, depositing a similar amount in each. This would result in a high, background-like, E_{ratio} value. To avoid this the algorithm could include a minimum distance between cells considered to be the two maxima, or try to identify minima between the two, this would greatly complicate a list-sorting approach.

The most complete, but resource-heavy, method might consist of fitting some functional form to the energies as a function of η and ϕ to extract the peak energies. This might work in software but is very computationally expensive, even if possible to implement in firmware it is likely not worth the latency it would require.

The desired solution is an algorithm for calculating an E_{ratio} -type variable that comes somewhere between these two options, more adaptable than the list-sorting approach and less resource-heavy than the peak fitting. The following sections explore such an algorithm: a baseline algorithm for finding secondary maxima in clusters in the GEP is established in Section 4.4.5.1; parameters of the algorithm are tuned using simulations in Sections 4.4.5.2, 4.4.5.3, and 4.4.5.4; and a summary of the results and recommended parameters, as well as additional adjustments that could be made with further studies, is given in Section 4.4.6.

4.4.5.1 Initial algorithm

Identifying the two highest energy cells is done in three stages: locating the seed, identifying candidate secondary maxima, and comparing results.

The GEP will receive a seed location from the eFEX identifying which SuperCell has the highest energy. The cells within this SuperCell are compared with one another to find which has the highest energy, this becomes the seed cell for the E_{ratio} algorithm.

The algorithm will then perform a stepwise search from the seed outwards to identify peaks in energy. On each step the energy gradient is calculated as $\Delta E = E_{\text{next}}^{\text{cell}} - E_{\text{prev}}^{\text{cell}}$, where $E_{\text{next}}^{\text{cell}}$ is the energy of the cell being stepped to, and $E_{\text{prev}}^{\text{cell}}$ is the energy of the cell being stepped from. From the first step ΔE should be negative, as the seed will have a higher energy than the surrounding cells, but on subsequent steps ΔE may become positive, marking that a minimum-energy point has been passed. If, after this, ΔE becomes negative again it indicates that the previous cell was a local maximum; in this case that cell is added to a list of candidate secondary maxima, and the search stops along this route. If the edge of the available range of cells is reached before ΔE turns positive then no candidate is saved. If the edge is reached after ΔE turns positive, but before it turns negative again, then the last cell in the range is taken to be the candidate.

This stepwise search is done in 6 different routes from the seed: one route where each step from the seed is in positive η , one in negative η , two where the first step is in positive ϕ before proceeding in positive or negative η , and two following the same pattern with the first step in negative ϕ . Figure 4.11 cells, alongside a schematic depicting the peak location strategy.

Once the stepwise search is complete, up to 6 candidate secondary maxima will have been identified. The candidate with the largest energy is taken as the secondary maximum and, with the seed as the maximum, E_{ratio} can be calculated using Equation 4.1.




Figure 4.11: Diagram showing the 6 different routes in which the E_{ratio} algorithm searches for secondary maxima (left) and how the algorithm identifies secondary maxima by tracking energy gradients along each step (right).

The performance of this baseline algorithm was investigated using simulations. Figure 4.12 shows the results, comparing the response in signal and background as a function of the calculated E_{ratio} value and the fraction of each that would pass a given E_{ratio} threshold. The background rejection as a function of signal efficiency is also shown, the baseline algorithm achieves a background rejection of 2.3 at 95% signal efficiency.

4.4.5.2 Peak size

The first parameter to investigate is the size of the area used to calculate energies. In the algorithm as described in Section 4.4.5.1, the energies used in comparisons and in the final E_{ratio} calculation are always the energies of a single cell. This could be modified by instead summing the energy of a cell with that of its neighbours in η to reduce sensitivity to small fluctuations. The number of cells summed is labelled the ‘peak size’, where the default algorithm would have a peak size of one. With a peak size greater than one the algorithm uses a ‘sliding window’ approach, so the step size is still a single cell despite the energy value coming from a larger area. For




Figure 4.12: Performance of baseline E_{ratio} algorithm on signal ($Z \rightarrow ee$) and background ($JZ0W$) clusters. Plots show (a) a histogram of calculated E_{ratio} values for each cluster, (b) the integral of (a) with a grey dashed line indicating the values at 95% signal efficiency, and (c) the background rejection of an E_{ratio} threshold corresponding to a given signal efficiency.




Figure 4.13: Diagram showing which cells contribute to the energy sum for the seed cell (bright yellow) for different peak sizes. The horizontal axis represents η and the vertical axis ϕ . In each case, the calculated energy would be the sum of the energies of the cells contained within the red box.

an odd numbered peak size the energy of a cell is added to that of its neighbours on each side. For an even numbered peak size, neighbours in positive η are preferred. Figure 4.13 shows how cells are included in the calculated energy.

Performance for the E_{ratio} algorithm was tested with peak size values from one to five. The results are shown in Figure 4.14. Comparing the benchmark value of background rejection at 95% signal efficiency, it is clear that a peak size of one (i.e. the same as the baseline algorithm) gives the best results, with performance degrading as more cells are added to the energy sum. This appears to be generally true for background rejection at all signal efficiencies. This suggests that the benefit of the fine granularity of each energy measurement outweighs the negative impact of any potential fluctuations that the increased peak size would smear out.

4.4.5.3 Exclusion region

Another alteration tested on the E_{ratio} algorithm is an ‘exclusion region’ around the seed, i.e. a number of cells close to the seed in which secondary maxima will not be searched for. An n -cell exclusion region means making the first step along any route n cells away from the seed in η . Since secondary maxima can be found as




Figure 4.14: Results for calculating E_{ratio} with different peak size options. Plots show background rejection as a function of signal efficiency for each peak size tested (left) and background rejection at 95% signal efficiency as a function of peak size (right).

soon as two steps have been taken from the seed cell (they cannot be found on the first step as it will always be a step down from the seed), this excludes all cells in an η range from $-n$ to n (in relative coordinates) from being considered secondary maxima. Comparatively, the baseline algorithm with no exclusion region can find secondary maxima anywhere but the four cells directly adjacent to the seed. Figure 4.15 highlights the effect of the exclusion region.

Performance for the E_{ratio} algorithm was tested with exclusion regions from between one and five cells, shown in Figure 4.16 alongside the baseline algorithm with no exclusion region. This time a clear increase in performance is visible compared to the initial form of the algorithm, with a one-cell exclusion region attaining a background rejection of 3.1 at 95% signal efficiency. For most signal efficiencies the one-cell exclusion still seems to perform best, though perhaps competing with a two-cell exclusion region for very high signal efficiencies. Since the only difference between no exclusion region and the one-cell case is that cells diagonally adjacent to the seed are excluded, these results suggest signal clusters frequently create secondary peaks on these diagonal; this could stem from incident particles falling close to the corner of a cell.




Figure 4.15: Diagram showing how the introduction of an exclusion region to the E_{ratio} algorithm prevents secondary maxima close to the seed from being selected. Red arrows mark each of the six paths traversed by the stepwise algorithm. Blue dots mark each step where the energy gradient is calculated. The shaded grey area shows cells that cannot be selected as a candidate secondary maximum, due to either being skipped over or being the first step from the seed.




Figure 4.16: Results for calculating E_{ratio} with different or no exclusion region definitions. Plots show background rejection as a function of signal efficiency for each tested exclusion region (left) and background rejection at 95% signal efficiency as a function of exclusion region size (right).




Figure 4.17: Plot of background rejection at 95% signal efficiency as a function of pseudorapidity, η , for E_{ratio} algorithms with different exclusion regions.

Given that cell widths vary significantly in different regions of the calorimeter, the performance for the E_{ratio} algorithm with different exclusion widths was also tested as a function of η . Figure 4.17 compares background rejection at 95% signal efficiency in several η regions. It is evident that the one-cell exclusion region performs best regardless of calorimeter geometry. The difference between one-cell and two-cell is much more drastic in the high- η endcap regions, here the strips are less granular so likely the larger exclusion regions are starting to miss real secondary peaks in background clusters.

4.4.5.4 Search limit

In the baseline E_{ratio} algorithm, the stepwise search for secondary maxima extends as far as the available data allows, in this case to the edge of the, conservatively large, stored cluster size. To minimise the amount of processing required by the algorithm, and potentially improve performance by reducing overlap with other clusters, a limit




Figure 4.18: Results for calculating E_{ratio} after varying the search limit parameter, given as a distance in η from the seed cell. Plots show background rejection as a function of signal efficiency for each tested search limit (left) and background rejection at 95% signal efficiency as a function of the search limit (right).

can be placed on the distance this search will traverse. Since the ϕ range of the search is already limited to one cell either side of the peak, this search limit is implemented as a maximum distance traversed in η . This distance is calculated in pseudorapidity units rather than number of cells to give a consistent response across calorimeter regions.

Performance for the E_{ratio} algorithm with different search limit values was tested on simulations, with the results presented in Figure 4.18. Distances in $\Delta\eta$ from 0.025 up to 0.15 were tested, with 0.15 being the width of the clusters and thus the limit in place in the baseline algorithm. While no performance gains are seen by reducing the search limit, there is a plateau in performance from $\Delta\eta > 0.1$. This means the required cluster size, and thus the amount of computation required, can be reduced without degrading performance of the algorithm.

4.4.6 Algorithm summary

Given the results presented in Sections 4.4.5.2-4.4.5.4, the most performant and resource-efficient algorithm for calculating E_{ratio} in the GEP is the baseline algorithm established in Section 4.4.5.1 with an exclusion region of one cell and a search limit of $\Delta\eta = 0.1$. No further improvement was found by varying the peak size. This algorithm achieves a background rejection of 3.1 for 95% signal efficiency.

The E_{ratio} algorithm presented here is functionally complete and serves as an option for e/γ discrimination in the GEP, but further improvements could be made with additional study. More parameters could be varied and tested for performance against simulations. One example is a threshold in the energy gradient between steps to allow a change in gradient to be identified, which might improve the response of the algorithm to noise or statistical fluctuations.

CHAPTER 5

Common Analysis Methods

5.1 Introduction

Two different analyses are presented in this thesis, although many methods, techniques, and properties of these analyses are shared. The analyses share a common final state, signal and background processes, selection for the $Z\gamma$ system, and more. The primary difference between the analyses being the phase space and treatment of the jets. This chapter discusses all of these shared analysis elements

5.2 Simulated event samples

The signal and background processes for both analyses are the same, the same simulated event samples can therefore be used for each.

Electroweak (EW) $Z\gamma jj$ production, the signal sample, is generated with MADGRAPH5_AMC@NLO 2.6.5 [7]. This sample is at Leading Order (LO) accuracy (order α_{EW}^4) with the NNPDF3.1 LO Parton Density Function (PDF) set [8]. Parton showering, hadronisation, and underlying event activity are added through PYTHIA 8.240 [9] (with “dipoleRecoil” enabled).

The dominant background process is Quantum Chromodynamics (QCD) production of $Z\gamma jj$. The nominal sample used for this process is produced with MADGRAPH5_AMC@NLO 2.3.3 [7] using the NNPDF3.0 Next-to-leading Order (NLO) PDF set [10], this includes all diagrams at order $\alpha_s^2\alpha_{EW}^2$ with up to two additional partons in the final state, where one parton may be at NLO.

Additional samples for QCD $Z\gamma jj$ are generated to evaluate uncertainties. A sample made with SHERPA 2.2.4 [11] at LO accuracy, with up to three additional parton emissions, is used to measure generator differences. The NNPDF3.0 Next-to-next-to-leading Order (NNLO) PDF set is used for this sample, in conjunction with a dedicated parton shower tuning developed by the SHERPA authors. Five more samples are used for evaluation of theoretical uncertainty. These are generated at particle level using SHERPA 2.2.10 [11] with the NNPDF3.0 NNLO PDF set. One of the five samples has a nominal value for the merging and resummation scales and the other four have an up or down variation for either.

The Z +jets background, in which a jet is misidentified as a photon, is estimated with a data-driven method. A Monte Carlo (MC) sample for this process is necessary to evaluate correlation between the regions used in this method, as discussed in Section 6.4. POWHEGBOX v1 [12, 5, 13] is used to generate this sample at NLO accuracy with the CT10 NLO PDF set [14]. PYTHIA 8.210 [9] is used for parton showering in this sample, with the AZNLO [15] set of tuned parameters.

A $t\bar{t}\gamma$ sample is generated at LO accuracy with MADGRAPH5_AMC@NLO and the NNPDF2.3 LO PDF set [16]. The WZ background has both a QCD and EW component; with samples from SHERPA 2.2.2 [11] at NLO, with NNPDF3.0 NNLO

PDF set, and MADGRAPH5_AMC@NLO 2.6.2 [7] at LO, with NNPDF3.0 LO PDF set, respectively.

Pileup (additional proton-proton interactions) is overlaid on simulated samples, generated with PYTHIA 8.186 using the A3 tune [17] and the NNPDF2.3 LO PDF set. Data is used to reweight these MC events to respect the mean number of interactions per bunch crossing from the corresponding data-taking period.

Once physics events and pileup are combined, samples are passed through a simulation of the ATLAS detector with GEANT4 [18, 19] and then processed with offline reconstruction in the same manner as data events. Additional scale factors and smearing are applied to more closely match data events.

CHAPTER 6

Search for vector-boson scattering production of a Z boson and a photon

6.1 Introduction

Vector-boson scattering provides a unique experimental signature, producing decay channels with excellent potential to probe rare Standard Model (SM) processes; the high selection efficiency achievable by exploiting kinematics of the Vector-boson Scattering (VBS) tag jets allows measurements to be conducted at lower cross-sections. The $Z\gamma jj$ final state can be interpreted as VBS $Z\gamma$ production, with the Z boson and photon resulting from a direct multiboson interaction. This provides a robust framework for studying these rare SM interactions, motivated previously in Chapter X. Feynman diagrams for VBS $Z\gamma$ production are represented in Figure 6.1, showing that Quadrilinear Gauge Coupling (QGC) or Trilinear Gauge Coupling (TGC) vertices are the only SM contributions at tree-level.




Figure 6.1: Feynman diagram for a $Z\gamma$ vector-boson scattering event (left). The black circle contains the multiboson interaction, which for a tree-level SM interaction will be one of the two shown (right).

In the archetypal VBS event, a quark from each of the two colliding protons radiates a boson. The two bosons interact to produce the EW component of the final state and the quarks, deflected from their original trajectories after boson emission, appear as jets in the detector. Since the initial quarks are typically very high energy, the angle through which the quarks are deflected is expected to be small. The final state jets, known as tag jets, would therefore be in the very forward regions of the detector, at opposite ends to one another, and also still carrying large amounts of energy. These kinematics are typically characterised by a large invariant mass of the dijet system (m_{jj}) and a large difference between the rapidities of the jets ($\Delta y(j, j)$).

VBS $Z\gamma$ production is one component of the more general EW production of $Z\gamma jj$. The VBS production modes are not gauge-invariantly separable from others, so a direct measurement of VBS $Z\gamma$ is not strictly possible. Instead, EW $Z\gamma jj$ production is measured with a selection designed to enhance the VBS component, matching the kinematics of the jets with the expected VBS signature. Figure 6.2 gives Feynman diagrams for some non-VBS production modes that contribute to the EW production mechanism.

To measure this EW $Z\gamma jj$ production, background processes with the same final




Figure 6.2: Example Feynman diagrams for non-VBS EW production of $Z\gamma jj$. In these instances the two final-state bosons do not originate from the same vertex.




Figure 6.3: Example Feynman diagrams for QCD production of $Z\gamma jj$.

state must be understood. The dominant background for this analysis comes from QCD $Z\gamma jj$ production. Figure 6.3 gives example Feynman diagrams for this QCD production, which differs from the EW mode as the strong force either provides the interaction between the two quarks or otherwise generates the final state jets. These strong interactions mean that the jets created must have compatible colours to facilitate further strong interactions: the jets are colour-connected. Additional interactions between the colour-connected jets are very probable and will affect the observed jet kinematics, allowing these events to be distinguished from VBS events.

6.2 Event selection

Additional requirements are placed on the transverse energies of these objects. The leading (highest energy) lepton must have $p_T > 30$ GeV, with the second lepton satisfying $p_T > 20$ GeV. Photons are required to have $p_T > 25$ GeV. The two highest energy jets must have $p_T > 50$ GeV.

For the dilepton system to be focused on $Z \rightarrow ll$ events, a cut of $m_{ll} > 40$ GeV is enforced, removing contributions from low-mass resonances. Events where the photon arises from final state radiation, emission from one of the leptons from the Z boson, are undesirable as the photon must exist at the same time as the Z boson for them to have come from the same vertex. A requirement on $m_{ll} + m_{ll\gamma} > 182$ GeV will suppress any events with a $Z \rightarrow ll\gamma$ decay producing the photon.

Imposing VBS-like kinematics on the jets further reduces the phase space. The dijet system should have a mass $m_{jj} > 150$ GeV, and a separation between the two jets of $|\Delta y_{jj}| > 1.0$. A veto is placed on ‘gap jets’, jets (reconstructed with $p_T > 25$ GeV) found in the rapidity region between the two VBS tag jets, is applied to exploit the difference between VBS jets and colour-connected QCD jets. A requirement on the centrality of the $ll\gamma$ system, relative to the jets, is also used to define the signal region. The centrality is defined by

$$\zeta(ll\gamma) = \left| \frac{y_{ll\gamma} - (y_{j1} + y_{j2})/2}{y_{j1} - y_{j2}} \right|,$$

separating the signal region ($\zeta(ll\gamma) < 0.4$) from the control region for the QCD background ($\zeta(ll\gamma) > 0.4$), explained in Section 6.4.

6.3 Particle-flow jet validation

A jet, being the name for a hadronic shower observed in the detector rather than an actual particle itself, can be defined in many different ways. The standard within ATLAS was to use ‘topo’ jets, formed from topological clusters of cells in the calorimeter [20]. More recently a move has been made to using particle-flow, or ‘PFlow’, jets. PFlow jets improve on the performance of topo jets by using tracks found in the tracking detector near to deposits in the calorimeter. This allows PFlow jets to benefit from advantages that the tracker has over the calorimeter: superior spacial resolution and improved energy resolution for low momentum particles [21].

Table 6.1: Yields and efficiencies after each jet cut, compared for both topo and PFlow jets. Starting from all EW $Z(\rightarrow ee)\gamma jj$ events that pass non-jet selection criteria.

Cut	Topo		PFlow		Difference
	Yield	Eff.	Yield	Eff.	
$N_j \geq 2$	51084	79.7%	51468	80.3%	+0.7%
$p_T^j > 50 \text{ GeV}, \eta_j < 4.5$	31362	61.4%	31604	61.4%	+0.8%
$\Delta R(l, j), \Delta R(\gamma, j) > 0.4$	31359	99.99%	31552	99.84%	+0.6%
$ \Delta\eta_{jj} > 1.0$	27127	86.5%	27293	86.5%	+0.6%
$ m_{jj} > 150 \text{ GeV}$	26752	98.6%	26885	98.5%	+0.5%

Whilst PFlow jets have been shown to improve on topo jets in general, they are not universally better in any scenario, each has its advantages. The most significant improvements for PFlow jets are in low p_T regimes, and PFlow can only function in the range of the tracker ($|\eta| < 2.5$). Jets in VBS events would not typically follow these criteria, they are expected to be very energetic and in the high- η forward regions of the detector. This study was therefore performed to validate that the new PFlow jets were still a good choice for this analysis in spite of this logical disconnect.

The procedure for comparing performance between these two jet collections is to investigate the difference in event yield from applying jet-based selection criteria with values calculated from each jet type. All non-jet selection criteria discussed in Section 6.2 are applied first, then jet-based selection criteria are applied and the resulting yields compared. For simplicity, this was investigated in the signal MC sample for $Z \rightarrow ee$ (but not $Z \rightarrow \mu\mu$) events.

Looking at the overall yield of events after each selection criterion shows that the two jet collections give very similar results, always equal to within 1%. Table 6.1 shows these yields for the signal MC sample.

Differences can be further scrutinised by looking at each individual event; most events should result in the same decision, pass or fail, regardless of the jet collection




Figure 6.4: Impact of four of the key jet cuts on the analysis yield when applied in both PFlow and topo. Events are divided into four categories based on whether or not they pass the PFlow cut and whether or not they pass the topo cut. Cuts are applied in the same order as presented in Table 6.1. Only events that passed the previous cut in both PFlow and topo are included in the results, to decorrelate the effects of each individual cut.

chosen. This checks that the similar yields aren't merely a coincidence, when in fact many events pass only one selection. Figure 6.4 shows these per-event differences in decisions.

For the vast majority of events, the two jet collections make the same selection, with less than 5% of events showing differences. These differences can be further scrutinised by investigating the distributions in the cut variables for these cases; if an event passes $m_{jj} > 150 \text{ GeV}$ with topo jets, but not with PFlow jets, the PFlow m_{jj} value should at least be close to 150 GeV.

Investigating the case where events pass selection with PFlow jets but fail with topo jets, results for the four key cut variables are shown in Figure 6.5. The distributions are, largely, as expected, with peaks showing most values are close to the cut

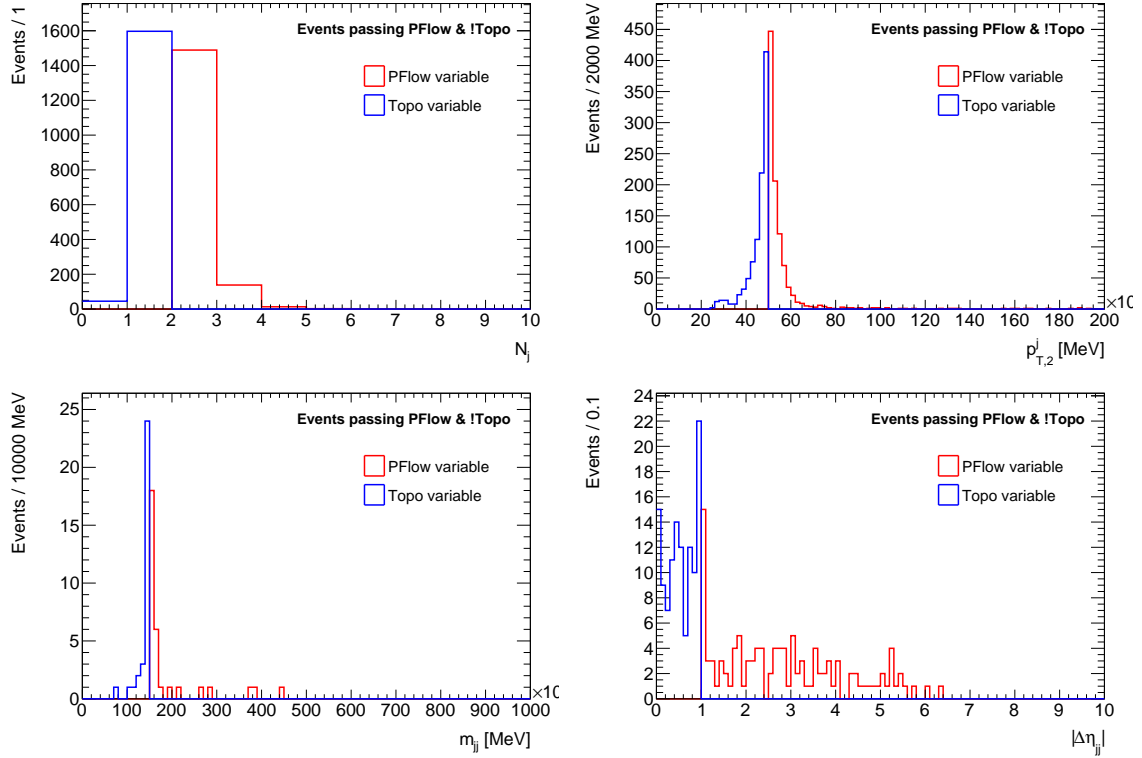


Figure 6.5: Distribution of events passing jet requirements for the PFlow collection but not the topo collection. In each case the cut is in the same distribution as the histogram being plotted. Four cuts are shown: $N_j > 2$ (top left), $p_T^{j,2}$ (p_T of second most energetic jet) > 50 GeV (top right), $m_{jj} > 150$ GeV (bottom left), $|\Delta\eta_{jj}| > 1$ (bottom right). Only events passing all prior cuts in both PFlow and topo are included. For the $p_T^{j,2}$ and m_{jj} plots the x -axis must be scaled by 10^3 .

boundaries. It is notable, however, that for the two variables using multiple jets, m_{jj} and $|\Delta\eta_{jj}|$, there are some strong outliers.

Events falling very far from the cut boundary in PFlow when the topo variable fails the cut seem to indicate a significant disagreement between kinematics in PFlow and topo jets. It is possible however that these outliers happen when the p_T ordering of jets varies between the two collection. Dijet variables are calculated using the two highest energy jets, so a small shift in p_T between the second and third jets could cause dijet variables to be calculated with a different jet pair and therefore give very different results.

This hypothesis can be tested by looking at the separation between jets used in each event for the two collections. The variable used to measure this is

$$\min \Sigma(\Delta R) = \min_{k_i \in \{(1,2), (2,1)\}} \sum_{i=1}^2 \Delta R(j_i^{\text{topo}}, j_{k_i}^{\text{PFlow}}), \quad (6.1)$$

i.e. the sum of the two ΔR values between topo and PFlow jets, for whichever pairing of the jets gives the lowest value of the sum. ΔR is the sum in quadrature of $\Delta\eta$ and $\Delta\phi$. Figure 6.6 shows the distribution of this variable for events passing m_{jj} in only PFlow. There are clearly two populations, separated at $\min \Sigma(\Delta R) \sim 0.5$. The lower population should contain events where the PFlow and topo jets are representing the same physical objects. Requiring $\min \Sigma(\Delta R) < 0.5$ on top of the existing selection gives the m_{jj} distribution shown in Figure 6.6, where now all remaining events are tightly distributed around the cut boundary.

This study, although limited in scope, serves to demonstrate that the key jet variables used in this analysis perform very similarly in topo jets and PFlow jets. Further studies could shed light on how systematic uncertainties differ between collections, and which choice would give the most precise result, but this was considered sufficient to update to the latest ATLAS recommendations.




Figure 6.6: Distribution of the $\min \Sigma(\Delta R)$ variable, defined in Equation 6.1, for events passing $m_{jj} > 150$ GeV in PFlow and not topo (left); and the m_{jj} distribution for these events but requiring $\min \Sigma(\Delta R) < 0.5$ (right).

6.4 Background estimation

The dominant background for this analysis, QCD $Z\gamma jj$ production, suffers from known mismodelling in areas with high dijet mass, which is precisely the region of interest here. Rather than using the MC template directly for estimating this background, the normalisation is corrected by comparing with data in a CR enriched in this background. The centrality variable is used to separate the SR ($\zeta(l\gamma) < 0.4$) from a QCD CR ($\zeta(l\gamma) > 0.4$). The CR is rich in the QCD background and has a very small fraction of signal events as the EW production mechanism peaks at low values of centrality. Figure 6.7 shows the centrality distribution for signal and backgrounds.

The normalisation correction is calculated by fitting a normalisation factor for the QCD background in the SR and CR simultaneously, allowing the overall normalisation to be adjusted according to data. The shape of the background is taken from MC, but data in the CR is used to validate the SR shape and constrain correlated uncertainties.

The second-largest background is from Z+jets events, and enters the final state when a jet is misidentified as a photon. This background includes Z+jets pro-




Figure 6.7: Centrality distribution for data and signal and background estimates pre-fit (before any data corrections to QCD $Z\gamma jj$). The dashed line marks the separation between the SR and CR. The uncertainty band is the quadrature sum of uncertainties from background estimation, MC statistics, and experimental systematics. Overflow events are included in the last bin. [4]

duced through both EW and QCD induced processes. The rate of particles being misidentified is not well modelled in MC, and so the shape and normalisation of this background must be corrected with a data-driven method.

The normalisation for the Z+jets background is calculated with the two-dimensional sideband, or ABCD, method. If the region in which the background is being estimated is region A, three more regions are defined by inverting isolation and identification criteria on the reconstructed photon. Inverting the photon isolation gives region B, inverting identification gives region C, and inverting both criteria gives region D. These three control regions are used to infer the amount of Z+jets background in the region of interest with the relationship

$$N_A^{Z+jets} = R \frac{N_B^{Z+jets} \times N_C^{Z+jets}}{N_D^{Z+jets}},$$

where N_X^{Z+jets} is the number of Z+jets events in the given region calculated by subtracting background and signal leakage from the data events i.e.

$$N_X^{Z+jets} = N_X^{data} - N_X^{bg} - c_X N_A^{sig,data}, \quad \text{for } X = B, C, D.$$

The correlation factor, R , is given by

$$R = \frac{N_A^{Z+jets} \times N_D^{Z+jets}}{N_B^{Z+jets} \times N_C^{Z+jets}},$$

where in this case each N_X^{Z+jets} is the event yield observed in Z+jets MC in this region. Also defined are signal leakage parameters, c_X , as

$$c_X = \frac{N_X^{sig}}{N_A^{sig}}, \quad \text{for } X = B, C, D,$$

calculated from QCD and EW $Z\gamma$ MC.

The shape of the Z+jets background is taken directly from a data control region. The control region should be very pure in Z+jets events, but also sufficiently high statistics. The chosen region is the anti-tight region, with no requirement on track

or calorimeter isolation. This is equivalent to regions C and D combined but without the track isolation requirement.

The background from $t\bar{t}\gamma$ events is estimated from MC and cross-checked in an $e\mu\gamma$ CR, which validates the use of a k -factor of 1.44 to scale the MC normalisation. $WZjj$ events make a minor contribution to the background, this is estimated solely from MC.

6.5 Systematic uncertainties

[This will largely be common between the two analyses, could be moved to analysis methods chapter]

This analysis considers uncertainties from data statistics and from systematics. Systematic uncertainties come from a variety of sources both theoretical and experimental.

Experimental systematic uncertainties cover uncertainties in energy scale and resolution of jets, photons, and electrons; momentum scale and resolution of muons; scale factors used to reproduce trigger, reconstruction, identification, and isolation efficiencies from data; suppression of pile-up jets; and flavour tagging.

Theoretical sources of systematic uncertainty come from the choice of scale, choice of PDF set, modelling, non-closure between MC generators, choice of parton showering and underlying event model, and EW-QCD interference. Uncertainty due to scale choice is calculated by varying the default values of renormalisation and factorisation scales in the nominal QCD $Z\gamma$ MC sample. Evaluating uncertainty in PDF set choice is done using the eigenvalues of the PDF set, for the signal and QCD $Z\gamma$ background. Uncertainty on modelling of the QCD $Z\gamma$ process comes from choice of merging (CKKW) and resummation (QSF) scale, these are calculated using the samples described in Section 5.2. For the QCD $Z\gamma$ background, the difference between the nominal and alternate samples is taken as a non-closure uncertainty. For




Figure 6.8: Plots of relative variation of yields as a function of dijet mass, m_{jj} , for EW (left) and QCD (right) production of $Z\gamma jj$. The variations shown are the largest groups of systematics in the SR. [4]

the EW signal, parton showering and underlying event uncertainties are calculated by comparing the default PYTHIA hadronisation to an alternate with HERWIG. The interference between EW and QCD $Z\gamma$ production is not included in either the signal or background, but instead taken as an additional uncertainty.

Further systematic uncertainties from limited statistics in MC samples and in data CRs are also considered.

These detailed systematics are used for both the signal and QCD $Z\gamma$ background. A selection of the largest groups of uncertainties are shown in the SR as a function of m_{jj} for both EW and QCD $Z\gamma jj$ production in Figure 6.8.

Uncertainties on the Z+jets background normalisation are calculated based on the two-dimensional sideband method, and total 35%. The $t\bar{t}\gamma$ and $WZjj$ backgrounds are both assigned flat uncertainties, of 15% and 20% respectively.

6.6 Jet flavour uncertainties

A significant source of systematic uncertainty in this analysis arises from uncertainty in the flavour of jets, in particular whether they are initiated by quarks or gluons. In order to minimise the impact of this uncertainty, the fraction of jets initiated by

gluons should be measured.

The gluon fraction is given by

$$f^{\text{gluon}} = \frac{N_j^{\text{gluon}}}{N_j^{\text{gluon}} + N_j^{\text{udsc}}},$$

where N_j^{gluon} and N_j^{udsc} are the number of jets initiated by gluons and by up/down/strange/charm quarks respectively. Measuring both f^{gluon} and its associated uncertainty as a function of jet pseudorapidity and transverse momentum provides the information needed to reduce the jet flavour uncertainty.

Uncertainty on f^{gluon} is from three sources: a modelling uncertainty calculated by finding the difference in f^{gluon} between two independent MC generators, statistical uncertainty from the size of the MC sample used, and an extra uncertainty to cover any differences in the value of f^{gluon} between regions. The third uncertainty component is necessary as, for technical reasons, only one f^{gluon} value could be provided for samples used to calculate yields in both the SR and the CR.

Calculation of f^{gluon} is performed for the QCD $Z\gamma$ sample only, as this is where the jet flavour uncertainty is largest.

Figure 6.9 shows the gluon fractions measured in the nominal and alternate MC samples for QCD $Z\gamma$, as well as calculations of each of the uncertainty components. The statistical uncertainty was found to be an order of magnitude smaller than the other components and so is neglected. The uncertainty to cover differences between regions is calculated by finding the largest difference, in each bin, between f^{gluon} in the inclusive region (SR+CR) and either of the two sub-regions. The overall uncertainty used is then the per-bin quadrature sum of the generator differences and this inter-region difference.




Figure 6.9: Gluon fractions and uncertainties as a function of jet pseudorapidity and transverse momentum. Plotted are gluon fraction in the nominal (a) and alternate (b) QCD $Z\gamma$ sample MC samples, the difference between gluon fractions in these two samples (c), the statistical uncertainty on gluon fractions in the alternate sample (d) (this was the larger of the statistical uncertainties), and the difference between gluon fraction in the SR+CR region and in the SR (e) and CR (f) regions.

6.7 Template fit

The signal strength, μ_{EW} , is used to parametrise the fiducial cross-section for the signal process, σ^{EW} , where

$$\mu_{\text{EW}} = \sigma_{\text{meas}}^{\text{EW}} / \sigma_{\text{exp}}^{\text{EW}},$$

i.e. the ratio of the measured cross-section to the SM expectation. This signal strength is extracted from the data with a maximum likelihood fit, performed on m_{jj} distributions in the SR and CR simultaneously. MC distributions for backgrounds and signal are used as templates, with normalisations for the signal and QCD $Z\gamma$ background allowed to float in the fit.

The bin size of the m_{jj} distribution used for the fit is chosen to optimise the statistical significance of data in each bin. Electron and muon channels are treated together, using the sum of events from both for all aspects of the fit. Likelihood functions are built for each bin of the m_{jj} distribution in both the SR and CR using a Poisson distribution, and these are all combined to form a single extended likelihood

Systematic uncertainties are included in the fit as Gaussian nuisance parameters. The effect of each uncertainty on the normalisation and shape of the m_{jj} distribution is considered individually and a pruning system, described in Section 6.8, is used to reduce the number of nuisance parameters needed.

The fit extracts the value of the signal strength for the signal process, μ_{EW} . From this a significance of the measurement, under the null hypothesis of $\mu_{\text{EW}} = 0$, can be extracted using the likelihood-ratio test. If the significance is greater than five standard deviations of a Gaussian, it is considered to be observed as an excess over the background-only hypothesis. The value of the signal strength can also be used to give a measurement of the fiducial cross-section of the process, given the of the SM expectation of this cross-section, calculated from the nominal MC:

$$\sigma_{\text{EW}}^{\text{exp}} = 4.73 \pm 0.01 \text{ (stat.)} \pm 0.15 \text{ (PDF)}_{-0.22}^{+0.23} \text{ (scale) fb.}$$

6.8 Pruning systematic uncertainties

To limit the number of nuisance parameters included in the fit, a system is developed to rank the impact of different systematics on both the shape and normalisation of the m_{jj} distribution. Only systematic uncertainties deemed to be significant are fully accounted for in the fit, and those with less impact are pruned.

The first test for a systematic uncertainty is how uniform its effect is across the dijet mass spectrum – this will indicate whether it will impact the shape of the m_{jj} distribution. If a systematic is determined to have a significant impact on shape, by criteria discussed below, then it is included in the fit with one nuisance parameter for each bin in m_{jj} , allowing it to modify the shape in the fitting process.

Any uncertainty not found to impact the shape should be assessed for how significant an impact it has on the overall normalisation of events. Systematic uncertainties with a large enough effect on the event yield will contribute one nuisance parameter to the fit, and have the ability to scale the overall normalisation. Any uncertainties with a smaller effect will be pruned, i.e. all pruned systematics will be added in quadrature as a single extra nuisance parameter to scale the overall normalisation in the fit.

6.8.1 Calculating statistical uncertainties

In order to determine whether the effect of any systematic uncertainty, on shape or overall yield, is significant, the statistical uncertainty on the value of the systematic uncertainty must be ascertained. This statistical uncertainty arises from the finite size of MC samples used to evaluate systematic uncertainties.

The value of a systematic uncertainty on an event yield is given by

$$\sigma_{\text{Norm}} = \frac{N_{\text{varied}} - N_{\text{nominal}}}{N_{\text{nominal}}}, \quad (6.2)$$

where N_{nominal} is the number of events in acceptance for a nominal MC sample and N_{varied} is the number of events after the systematic variation has been applied. Each of N_{varied} and N_{nominal} has a statistical uncertainty. However, due to the fact that these variables are measuring the same set of events under different conditions, the two yields are highly correlated. The correlation is not known a priori, and so the uncertainty on σ_{Norm} cannot be calculated through error propagation.

The bootstrap method[22, 23] is instead used in order to determine statistical uncertainties while preserving correlations. This method relies on resampling the event set to create replica sets of events of the same size, with some events duplicated and some omitted. Calculating σ_{Norm} in each replica set gives a distribution of results for which the standard deviation represents the statistical uncertainty on σ_{Norm} .

6.8.2 Determining shape impact

For a systematic uncertainty that has no impact on m_{jj} shape, it would be expected that the resulting variation would be uniform across the m_{jj} distribution. This is tested by calculating the fractional difference in yield, σ_{Norm} , and its associated statistical uncertainty in bins of m_{jj} . A chi-squared test from fitting a zeroth order polynomial to these values provides a test statistic which should be distributed as $\chi^2(3)$ (4 bins minus 1 parameter for 3 degrees of freedom) under the null hypothesis of no shape impact. A significant shape uncertainty is therefore anything that deviates from this null hypothesis above a certain threshold. An example m_{jj} distribution and fit is shown in Figure 6.10.

The threshold chosen is a p -value of 0.05, i.e. chi-squared values sufficiently high that there is at most a 5% chance that the deviation arose from statistical fluctuation. Uncertainties with a p -value below this threshold have the full shape treatment in the fit, with per-bin nuisance parameters. Figure 6.11 shows the results of the chi-squared test for the largest experimental systematics in the SR.




Figure 6.10: Binned m_{jj} distribution of measured systematic uncertainty, σ_{Norm} , for one systematic variation as labelled on the plot. The dashed line shows the zeroth-order polynomial fit attempted, with the indicated χ^2 value demonstrating this is clearly a poor assumption and this systematic does contribute an uncertainty on the m_{jj} shape.




Figure 6.11: χ^2 values, representing the impact each systematic has on shape, for all experimental systematic variations in the EW (top) and QCD (bottom) samples in the SR. Only sources with χ^2 above 7 for either the up or down variation are shown. The top axis gives the probabilities of uncertainties arising fluctuations under the null hypothesis. The largest uncertainties extend beyond the range of the x -axis.

6.8.3 Determining overall yield impact

Any systematic uncertainty determined to not impact the shape of the m_{jj} distribution can of course still affect the overall yield of events, and therefore require sufficient treatment in the fit. This could be done by assigning a single nuisance parameter to every remaining uncertainty, as all will have an effect on some scale. To reduce the number of nuisance parameters required however, the less significant uncertainties are pruned.

In this case significance is simply determined by whether or not a systematic uncertainty is consistent with zero. Taking the value of the systematic, per Equation 6.2, and its statistical uncertainty or standard deviation, if the value is within one standard deviation of zero it is considered consistent with being zero.

All systematic uncertainties not passing the shape significance test but not consistent with zero have a dedicated nuisance parameter in the fit. All remaining uncertainties are pruned. The relative change in yield from each of the largest experimental systematic uncertainties is shown in Figure 6.12.

6.9 Results

Fitted m_{jj} distributions are shown in Figure 6.13. The measured signal strength in the fit is

$$\begin{aligned}\mu_{EW} &= 0.95^{+0.14}_{-0.13} \\ &= 0.95 \pm 0.08 \text{ (stat.)} \pm 0.11 \text{ (syst.)}.\end{aligned}$$

This corresponds to an observed(expected) significance of 10(11) standard deviations, and is the first observation of this process by the ATLAS collaboration.

The fiducial cross-section of the EW production of $Z\gamma jj$ in this VBS-like phase space




Figure 6.12: Relative change in yield, σ_{Norm} , for all experimental systematic uncertainties in the SR for the EW (top) and QCD (bottom) samples. Any with a value below 5×10^{-4} are omitted. The black bars represent the statistical uncertainty on the value.

is therefore measured as

$$\sigma_{\text{EW}} = 4.49 \pm 0.40 \text{ (stat.)} \pm 0.42 \text{ (syst.) fb.}$$




Figure 6.13: Post-fit distributions of dijet mass, m_{jj} , in the SR (top) and CR (bottom). The uncertainty band is the combination of all uncertainties, taken from the fit. Overflow events are included in the last bin. [4]

CHAPTER 7

Search for triboson production of $VZ\gamma$ through its semi-leptonic decay mode

7.1 Introduction

Triboson production of a Z boson; a photon; and an additional, hadronically decaying, vector boson is another production mechanism for the $Z\gamma jj$ final state. This additional boson, denoted as a ‘V’ boson, can be a W or a Z boson. Measuring this $VZ\gamma$ triboson process thus constitutes an inclusive measurement of both $WZ\gamma$ and $ZZ\gamma$ triboson production.

Figure 7.1 gives tree-level Feynman diagrams for $WZ\gamma$ and $ZZ\gamma$ production. Notably, $WZ\gamma$ production is sensitive to the same QGC and TGC vertices as VBS $Z\gamma$ production. $ZZ\gamma$ is more limited by the lack of neutral QGCs and TGCs in the SM, but is therefore more sensitive to potential anomalous couplings. The reliance




Figure 7.1: A selection of SM production mechanisms for the $VZ\gamma$ triboson final state, depicted in Feynman diagrams.

on these rare multiboson interactions means that the triboson process has a low cross-section, similar to that of VBS $Z\gamma$. Without the distinct VBS jet signature to select on, measuring this low cross-section process is challenging.

The two jets, here a product of a boson decay, have kinematic properties that distinguish them from background events, notably: a dijet mass peaked around the W/Z masses, small rapidity separation between the two jets, and an angular distribution consistent with boson decay products. Additionally, more subtle differences are also present in many other variables. Selecting W/Z boson decays to leptons is relatively straightforward, but the more limited jet resolution and more dominant jet background makes doing this in the hadronic decay channel more difficult. This analysis employs machine-learning techniques to interpret this complicated phase space; pushing sensitivity beyond what is achievable with a traditional cut-based analysis.

Despite the differing jet phase space, the backgrounds for this analysis are the same as for the VBS $Z\gamma$ analysis. QCD $Z\gamma jj$ production is the dominant competing pro-

cess; the key difference between this background and the signal is the kinematics of the jets, as both have a real Z boson and photon produced. The QCD $Z\gamma$ background has a yield X times higher than the signal after applying selection on the Z boson and photon. This illustrates the need for effective jet selection to manage this background.

The goal of this analysis is to measure the fiducial cross-section of this rare process in order to compare it to the standard model expectation. This measurement is extracted from a template fit to the output of the machine-learning model. If the observed significance is sufficient, this will provide evidence for or an observation of this process. If the significance does not meet these thresholds, the measurement will be used to place limits on the rate of this process. These limits can be used to constrain theories that might enhance the cross-section of this process.

7.2 Discriminating against QCD $Z\gamma$ production

The biggest challenge in this analysis is managing the dominant background, QCD $Z\gamma jj$ production. Like the signal process, this background has a real Z boson and photon. The difference is the origin of the jets, here not from a boson decay but more likely radiated from the initial or final state. Identifying and exploiting the differences in jet kinematics between this background and the signal is therefore key to maximising the sensitivity of the measurement. This section is dedicated to discussing this problem; thus here the word signal is used to refer to EW $VZ\gamma$ production and background refers solely to QCD $Z\gamma jj$ production.

There are a small number of kinematic distributions which exhibit a large difference between signal and background that could be exploited effectively by a cut; the dijet mass, m_{jj} , being an obvious example as for the signal it peaks around the W/Z boson mass but for the background resembles a continuum. For a great many more variables however, the differences are more subtle. Whilst there is frequently an obvious difference in shape between signal and background, there is no obvious set

of cuts that would create a signal-rich region. Figure 7.2 shows some distributions with the largest signal-background discrepancies.

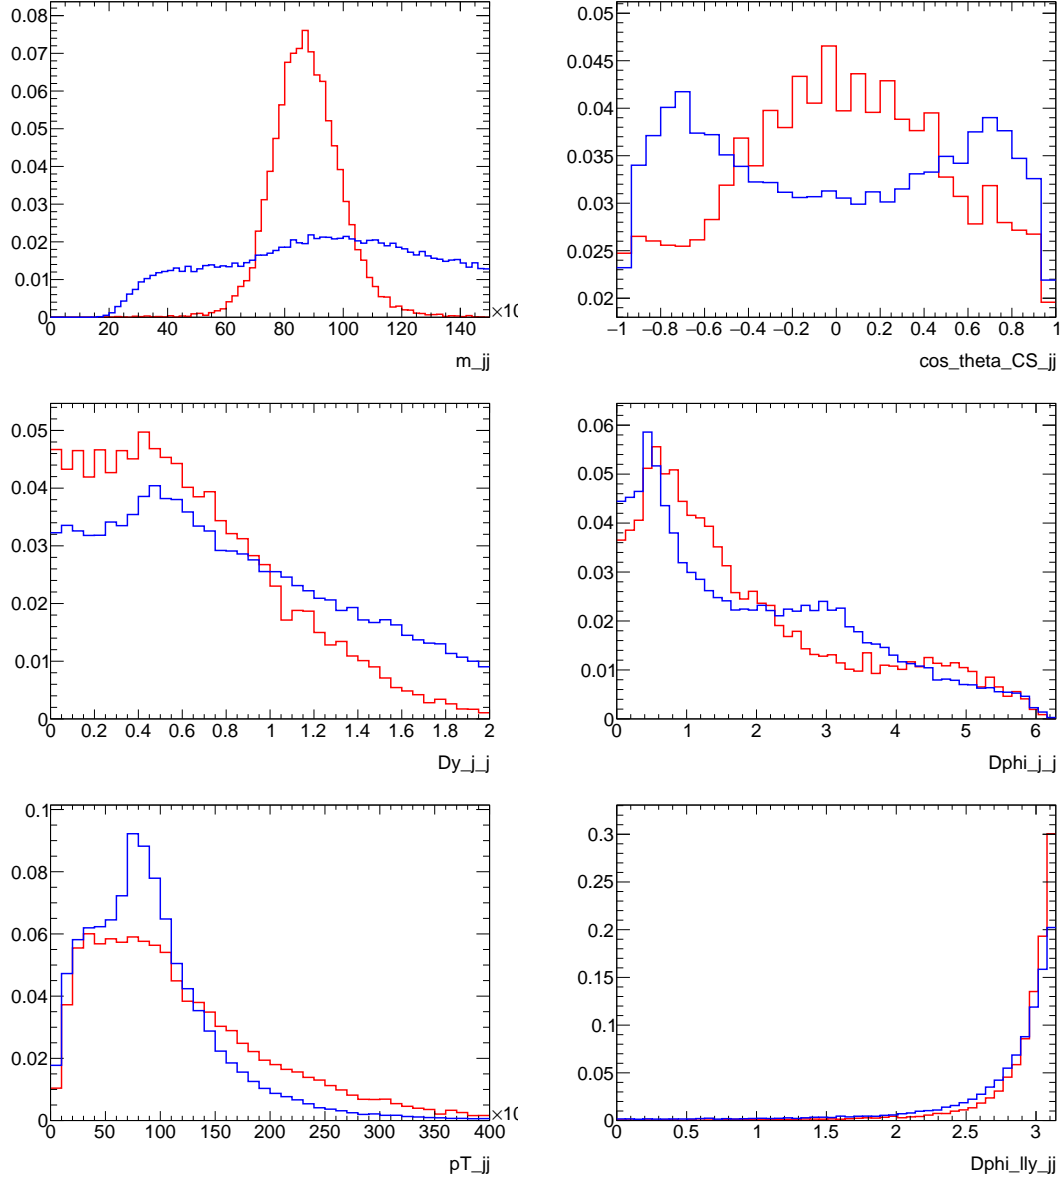


Figure 7.2: Kinematic distributions, comparing EW $VZ\gamma$ production (red) to QCD $Z\gamma\gamma$ production (blue). [Placeholder plots, to be updated]

CHAPTER 8

Conclusions

REFERENCES

- [1] I. Brawn. *L1Calo Overview, Status, Installation & Commissioning* (Oct 2019). <https://indico.cern.ch/event/829769/contributions/3572289>. [ATLAS Internal].
- [2] The ATLAS Collaboration. *Technical Design Report for the Phase-I Upgrade of the ATLAS TDAQ System*. Technical Report ATLAS-TDR-023-2013 (2013). <http://cds.cern.ch/record/1602235/>.
- [3] The ATLAS Collaboration. *Technical Design Report for the Phase-II Upgrade of the ATLAS TDAQ System*. Technical Report ATLAS-TDR-029, CERN, Geneva (2017). <https://cds.cern.ch/record/2285584>.
- [4] The ATLAS Collaboration. *Measurement of the cross-sections of the electroweak and total production of a $Z\gamma$ pair in association with two jets in pp collisions at $\sqrt{s} = 13$ TeV with the ATLAS detector* (Jul 2023). <http://cds.cern.ch/record/2779171><http://arxiv.org/abs/2305.19142>. 2305.19142.
- [5] S. Frixione, P. Nason, and C. Oleari. *Matching NLO QCD computations with parton shower simulations: The POWHEG method*. J. High Energy Phys. 2007(11) (Nov 2007) p. 070. ISSN 11266708. <http://dx.doi.org/10.1088/1126-6708/2007/11/070>. 0709.2092.
- [6] T. Sjöstrand, S. Mrenna, and P. Skands. *A brief introduction to PYTHIA 8.1*. Comput. Phys. Commun. 178(11) (Jun 2008) pp. 852–867. ISSN 00104655. <http://dx.doi.org/10.1016/j.cpc.2008.01.036>. 0710.3820.
- [7] J. Alwall et al. *The automated computation of tree-level and next-to-leading order differential cross sections, and their matching to parton shower simulations*. J. High Energy Phys. 2014 20147 2014(7) (Jul 2014) pp. 1–157. ISSN 1029-8479. [http://dx.doi.org/10.1007/JHEP07\(2014\)079](http://dx.doi.org/10.1007/JHEP07(2014)079). 1405.0301.

-
- [8] R. D. Ball et al. *Parton distributions from high-precision collider data*. Eur. Phys. J. C 2017 7710 77(10) (Oct 2017) pp. 1–75. ISSN 1434-6052. <http://dx.doi.org/10.1140/EPJC/S10052-017-5199-5>.
- [9] T. Sjöstrand et al. *An introduction to PYTHIA 8.2*. Comput. Phys. Commun. 191(1) (Jun 2015) pp. 159–177. ISSN 0010-4655. <http://dx.doi.org/10.1016/J.CPC.2015.01.024>. 1410.3012.
- [10] R. D. Ball et al. *Parton distributions for the LHC run II*. J. High Energy Phys. 2015 20154 2015(4) (Apr 2015) pp. 1–148. ISSN 1029-8479. [http://dx.doi.org/10.1007/JHEP04\(2015\)040](http://dx.doi.org/10.1007/JHEP04(2015)040). 1410.8849.
- [11] E. Bothmann et al. *Event generation with Sherpa 2.2*. SciPost Phys. 7(3) (Sep 2019) p. 034. ISSN 25424653. <http://dx.doi.org/10.21468/SCIPOSTPHYS.7.3.034/PDF>. 1905.09127.
- [12] P. Nason. *A new method for combining NLO QCD with shower Monte Carlo algorithms*. J. High Energy Phys. 8(11) (Dec 2004) pp. 1097–1124. ISSN 10298479. <http://dx.doi.org/10.1088/1126-6708/2004/11/040>. 0409146.
- [13] S. Alioli et al. *A general framework for implementing NLO calculations in shower Monte Carlo programs: The POWHEG BOX*. J. High Energy Phys. 2010(6) (Jun 2010) pp. 1–58. ISSN 10298479. [http://dx.doi.org/10.1007/JHEP06\(2010\)043](http://dx.doi.org/10.1007/JHEP06(2010)043). 1002.2581.
- [14] H. L. Lai et al. *New parton distributions for collider physics*. Phys. Rev. D - Part. Fields, Gravit. Cosmol. 82(7) (Oct 2010) p. 074,024. ISSN 15507998. <http://dx.doi.org/10.1103/PhysRevD.82.074024>. 1007.2241.
- [15] The ATLAS Collaboration. *Measurement of the Z/γ^* boson transverse momentum distribution in pp collisions at $\sqrt{s} = 7$ TeV with the ATLAS detector*. J. High Energy Phys. 2014(9) (Sep 2014) pp. 1–47. ISSN 10298479. [http://dx.doi.org/10.1007/JHEP09\(2014\)145/METRICS](http://dx.doi.org/10.1007/JHEP09(2014)145/METRICS). 1406.3660.
- [16] R. D. Ball et al. *Parton distributions with LHC data*. Nucl. Phys. B 867(2) (Feb 2013) pp. 244–289. ISSN 05503213. <http://dx.doi.org/10.1016/j.nuclphysb.2012.10.003>. 1207.1303.
- [17] The ATLAS Collaboration. *The Pythia8 A3 tune description of ATLAS minimum bias and inelastic measurements incorporating the Donnachie-Landshoff diffractive model*. ATL-PHYS-PUB-2016-017 (Aug 2016). <https://cds.cern.ch/record/2206965https://inspirehep.net/record/1477266>.
- [18] The ATLAS Collaboration. *The ATLAS Simulation Infrastructure*. Eur. Phys. J. C 70(3) (Sep 2010) pp. 823–874. ISSN 14346052. <http://dx.doi.org/10.1140/epjc/s10052-010-1429-9>. 1005.4568.
- [19] S. Agostinelli et al. *GEANT4 - A simulation toolkit*. Nucl. Instruments Methods Phys. Res. Sect. A Accel. Spectrometers, Detect. Assoc. Equip. 506(3)

- (Jul 2003) pp. 250–303. ISSN 01689002. [http://dx.doi.org/10.1016/S0168-9002\(03\)01368-8](http://dx.doi.org/10.1016/S0168-9002(03)01368-8).
- [20] The ATLAS Collaboration. *Topological cell clustering in the ATLAS calorimeters and its performance in LHC Run 1*. Eur. Phys. J. C 77(7) (2017) p. 490. ISSN 1434-6052. <http://dx.doi.org/10.1140/epjc/s10052-017-5004-5>.
- [21] The ATLAS Collaboration. *Jet reconstruction and performance using particle flow with the ATLAS Detector*. Eur. Phys. J. C 77(7) (2017) p. 466. ISSN 1434-6052. <http://dx.doi.org/10.1140/epjc/s10052-017-5031-2>.
- [22] B. Efron. *Computers and the Theory of Statistics: Thinking the Unthinkable*. SIAM Rev. 21(4) (Oct 1979) pp. 460–480. ISSN 0036-1445. <http://dx.doi.org/10.1137/1021092>.
- [23] B. Efron. *Better Bootstrap Confidence Intervals*. J. Am. Stat. Assoc. 82(397) (Mar 1987) p. 171. ISSN 01621459. <http://dx.doi.org/10.2307/2289144>.

APPENDIX A

FIRST APPENDIX
